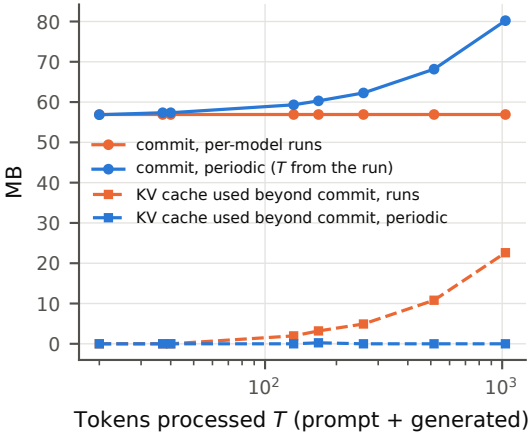
llama.cpp, SmolLM2-135M (glibc commits 405 MB)
MB
Tokens processed T (prompt + generated)
commit, per-model runs
commit, periodic (T from the run)
KV cache used beyond commit, runs
KV cache used beyond commit, periodic

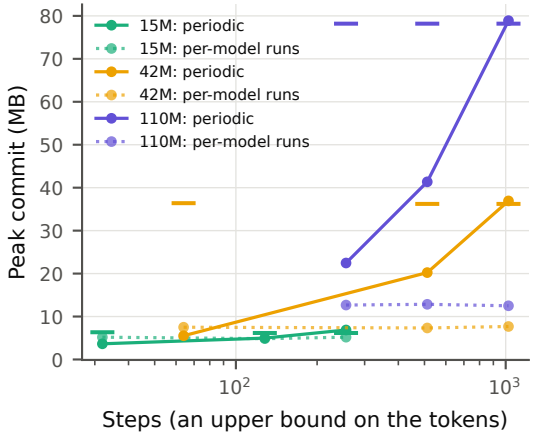
llama2.c (bars: glibc)
Peak commit (MB)
Steps (an upper bound on the tokens)
15M: periodic
15M: per-model runs
42M: periodic
42M: per-model runs
110M: periodic
110M: per-model runs