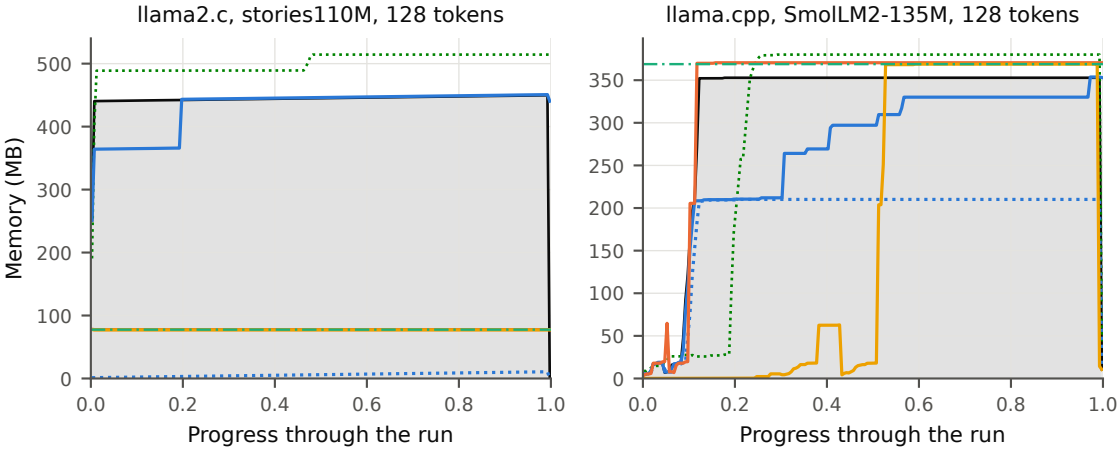
llama2.c, stories110M, 128 tokens
llama.cpp, SmolLM2-135M, 128 tokens
Memory (MB)
Progress through the run
Progress through the run
True footprint (bytes touched)
Resident set (RSS)
MemPrint windowed 5%
MemPrint pages only
Massif curve
heaptrack curve
DHAT peak